

Structured Concurrent Programming

William Cook, Jayadev Misra,
David Kitchin, Adrian Quark,
Andrew Matsuoka, John Thywissen

Department of Computer Science
University of Texas at Austin

<http://orc.csres.utexas.edu>

Outline

Overview

Orc Notation

Examples

Laws

A Time-Based Algorithm

Status of Research

Structured Concurrent Programming

- **Structured Sequential Programming:** Dijkstra circa 1968
Fundamental Question: Component Integration.
- **Structured Concurrent Programming:** Component Integration
 - Concurrency combinators that promote component integration
 - Paradigms for constructing concurrent and distributed programs
 - Orchestration

Some Typical Applications

- **Account management in a bank** (Business process management):
Workflow lasting over several months
Security, Failure, Long-lived Data
- **Extended 911**:
Using humans as components
Components join and leave
Real-time response
- **Network simulation**:
Experiments with differing traffic and failure modes
Animation
- **Managing a city**: (A proposal to EU)
Components integrated dynamically
The scope of software is nebulous

Some Typical Applications, contd.

- Matrix computation in a multi-core machine
- Map-Reduce using a server farm
- Concurrency management in database access
- Thread management in an operating system
- Mashups (Internet Scripting)

Internet Scripting

- Contact two airlines simultaneously for price quotes.
- Buy a ticket if the quote is at most \$300.
- Buy the cheapest ticket if both quotes are above \$300.
- Buy a ticket if the other airline does not give a timely quote.
- Notify client if neither airline provides a timely quote.

-

Orchestrating Components (services)

Acquire data from services.

Calculate with these data.

Invoke yet other services with the results.

Additionally

Invoke multiple services simultaneously for failure tolerance.

Repeatedly poll a service.

Ask a service to notify the user when it acquires the appropriate data.

Download a service and invoke it locally.

Have a service call another service on behalf of the user.

...

Orc, an Orchestration Theory

- **Site**: Basic service (component).
- Concurrency **combinators** for integrating sites.
- Theory includes nothing other than the combinators.

No notion of data type, thread, process, channel,
synchronization, parallelism ...

New concepts are programmed using the combinators.

Sites

- **External Services:** Google spell checker, Google Search, MySpace, CNN, Discovery ...
- **Giant Components:** Linux, Homeland Security Database ...
- **Any Java Class instance**
- **Library sites**
 - `+ - * && || ...`
 - `println, random, Prompt, Email ...`
 - `Timer`
 - `Database, Semaphore, Channel ...`
 - `Sites that create sites: MakeDatabase, MakeSemaphore, MakeChannel ...`
 - ...

Overview of Orc

- Orc program has
 - a **goal** expression,
 - a set of definitions.
- The goal expression is executed. Its execution
 - calls **sites**,
 - publishes **values**.
- Orc is simple
 - Orc has only 3 combinators.
 - Can handle time-outs, priorities, failures, synchronizations, ...
 - Implementation allows writing simple expressions.
 $2 + 3$ is translated to site call `Add(2,3)`.

Structure of Orc Expression

- **Simple**: just a site call, $CNN(d)$
Publishes the value returned by the site.

- **Composition** of two Orc expressions:

do f and g in parallel	$f \mid g$	Symmetric composition
for all x from f do g	$f >x> g$	Sequential composition
for some x from g do f	$f <x< g$	Pruning

Structure of Orc Expression

- **Simple**: just a site call, $CNN(d)$
Publishes the value returned by the site.

- **Composition** of two Orc expressions:

do f and g in parallel	$f \mid g$	Symmetric composition
for all x from f do g	$f >x> g$	Sequential composition
for some x from g do f	$f <x< g$	Pruning

Structure of Orc Expression

- **Simple**: just a site call, $CNN(d)$
Publishes the value returned by the site.

- **Composition** of two Orc expressions:

do f and g in parallel	$f \mid g$	Symmetric composition
for all x from f do g	$f > x > g$	Sequential composition
for some x from g do f	$f < x < g$	Pruning

Structure of Orc Expression

- **Simple**: just a site call, $CNN(d)$
Publishes the value returned by the site.

- **Composition** of two Orc expressions:

do f and g in parallel	$f \mid g$	Symmetric composition
for all x from f do g	$f >x> g$	Sequential composition
for some x from g do f	$f <x< g$	Pruning

Symmetric composition: $f \mid g$

- Evaluate f and g independently.
- Publish all values from both.
- No direct communication or interaction between f and g .
They can communicate only through sites.

Examples

- $CNN(d) \mid BBC(d)$: calls both CNN and BBC simultaneously.
Publishes values returned by both sites. (0, 1 or 2 values)
- $WebServer() \mid MailServer() \mid LinuxServer()$
May not publish any value.

Sequential composition: $f \text{ > } x \text{ > } g$

For all values published by f do g .

Publish only the values from g .

- $CNN(d) \text{ > } x \text{ > } Email(address, x)$
 - Call $CNN(d)$.
 - Bind result (if any) to x .
 - Call $Email(address, x)$.
 - Publish the value, if any, returned by $Email$.
- $(CNN(d) \mid BBC(d)) \text{ > } x \text{ > } Email(address, x)$
 - May call $Email$ twice.
 - Publishes up to two values from $Email$.

Schematic of Sequential composition




Figure: Schematic of $f >x> g$

Pruning: $(f \text{ } <x< \text{ } g)$

For some value published by g do f .

Publish only the values from f .

- Evaluate f and g in parallel.
 - Site calls that need x are suspended.
 - Other site calls proceed.
 - see $(M() \mid N(x)) \text{ } <x< \text{ } g$
- When g returns a (first) value:
 - Assign it to x .
 - Terminate g .
 - Resume suspended calls.
- Values published by f are the values of $(f \text{ } <x< \text{ } g)$.

Example of Pruning

$Email(address, x) \text{ } <x< (CNN(d) \mid BBC(d))$

Binds x to the first value from $CNN(d) \mid BBC(d)$.
Sends at most one email.

-

Some Fundamental Sites

- $if(b)$: boolean b ,
returns a **signal** if b is true; remains **silent** if b is false.
- $Rtimer(t)$: integer t , $t \geq 0$, returns a signal t time units later.
- $stop$: never responds. Same as $if(false)$.
- $signal()$ returns a signal immediately. Same as $if(true)$.

Centralized Execution Model

- An expression is evaluated on a single machine (**client**).
- Client communicates with sites by messages.

Time-out

Publish M 's response if it arrives before time t ,
Otherwise, publish 0 .

$$z \prec z \prec (M() \mid (Rtimer(t) \gg 0))$$

-

Fork-join parallelism

Call M and N in parallel.

Return their values as a tuple after both respond.

$$((u, v) \text{ } \textcolor{red}{<u}< M()) \text{ } \textcolor{red}{<v}< N())$$

-

Expression Definition

def *MailOnce*(*a*) =
 Email(*a*, *m*) $<_m$ (*CNN*(*d*) | *BBC*(*d*))

def *MailLoop*(*a*, *d*) =
 MailOnce(*a*) \gg *Rtimer*(*d*) \gg *MailLoop*(*a*, *d*)

- Expression is called like a procedure.
It may publish many values. *MailLoop* does not publish.
- Site calls are strict; expression calls non-strict.

Expression Definition

- output n signals -

def $signals(n) = if(n > 0) \gg (signal \mid signals(n - 1))$

- Publish a signal at every time unit.-


def $metronome() = signal \mid (Rtimer(1) \gg metronome())$

- Publish a signal every t time units.-

def $tmetronome(t) = signal \mid (Rtimer(t) \gg tmetronome(t))$

- Publish natural numbers from i every t time units.-

def $gen(i, t) = i \mid Rtimer(t) \gg gen(i + 1, t)$



Recursive definition with time-out

Call a list of sites.

Count the number of responses received within 10 time units.

def *tally*([]) = 0

def *tally*(*M* : *MS*) = *u* + *v*

$\langle u \rangle \leftarrow (M() \ggg 1) \mid (Rtimer(10) \ggg 0)$

$\langle v \rangle \leftarrow tally(MS)$

or, in the current language,

def *tally*([]) = 0

def *tally*(*M* : *MS*) = (*M*() \ggg 1 \mid *Rtimer*(10) \ggg 0) + *tally*(*MS*)

Barrier Synchronization in $M \gg f \mid N \gg g$

f and g start only after **both** M and N complete.

Rendezvous of CSP or CCS; M and N are complementary actions.

$$\begin{aligned} & ((u, v) \\ & \quad <u < M() \\ & \quad <v < N()) \\ & \gg (f \mid g) \end{aligned}$$

Priority

- Publish N 's response asap, but no earlier than 1 unit from now.
Apply fork-join between $Rtimer(1)$ and N .

$$\text{def } Delay() = (Rtimer(1) \gg u) <u < N()$$

- Call M , N together.
If M responds within one unit, publish its response.
Else, publish the first response.

$$x <x < (M() \mid Delay())$$

Interrupt f

Evaluation of f can not be directly interrupted.

Introduce two sites:

- *Interrupt.set*: to interrupt f
- *Interrupt.get*: responds after *Interrupt.set* has been called.

Instead of f , evaluate

$$z \leq z \leq (f \mid \text{Interrupt.get}())$$

Parallel or

Sites M and N return booleans. Compute their **parallel or**.

$$\begin{aligned} & \text{if}(x) \gg \text{true} \mid \text{if}(y) \gg \text{true} \mid (x||y) \\ & \quad <x< M() \\ & \quad <y< N() \end{aligned}$$

To return just one value:

$$\begin{aligned} & z \\ & \quad <z< \text{if}(x) \gg \text{true} \mid \text{if}(y) \gg \text{true} \mid (x||y) \\ & \quad <x< M() \\ & \quad <y< N() \end{aligned}$$

Airline quotes: Application of Parallel or

Contact airlines A and B .

Return any quote if it is below \$300 as soon as it is available,
otherwise return the minimum quote.

$threshold(x)$ returns x if $x < 300$; silent otherwise.

$Min(x, y)$ returns the minimum of x and y .

z

$\langle z \rangle < threshold(x) \mid threshold(y) \mid Min(x, y)$

$\langle x \rangle < A()$

$\langle y \rangle < B()$

Backtracking: Eight queens




Figure: Backtrack Search for Eight queens

Eight queens; contd.

def *extend*(*z*, 1) = *valid*(0:*z*) | *valid*(1:*z*) | ... | *valid*(7:*z*)

def *extend*(*z*, *n*) = *extend*(*z*, 1) *>y>* *extend*(*y*, *n* - 1)

- *z*: partial placement of queens (list of values from 0..7)
- *extend*(*z*, *n*) publishes **all** valid extensions of *z* with *n* additional queens.
- *valid*(*z*) returns *z* if *z* is valid; silent otherwise.
- Solve the original problem by calling *extend*([], 8).

Processes

- Processes typically communicate via channels.
- For channel c , treat $c.put$ and $c.get$ as site calls.
- In our examples, $c.get$ is blocking and $c.put$ is non-blocking.
- Other kinds of channels can be programmed as sites.

Typical Iterative Process

Forever: Read x from channel c , compute with x , output result on e :

def $P(c, e) = c.get() \rightarrow x \rightarrow \text{Compute}(x) \rightarrow y \rightarrow e.put(y) \gg P(c, e)$




Figure: Iterative Process

Process Network

Process (network) to read from both c and d and write on e :

$$\text{def } \text{Net}(c, d, e) = P(c, e) \mid P(d, e)$$





Figure: Network of Iterative Processes

Workload Balancing

Read from c , assign work randomly to one of the processes.

```
def bal(c, c', d') =
  c.get() >x> random(2) >t>
  ( if(t = 0) >> c'.put(x)
    | if(t = 1) >> d'.put(x)
  ) >> bal(c, c', d')
```

```
def WorkBal(c, e) =
  Buffer() >c'> Buffer() >d'>
  (bal(c, c', d') | Net(c', d', e))
```



WorkBal(c,d,e)

Figure: Workload Balancing in a network of Processes

Network of Services: Insurance Company

```
def insurance() = apply() | join() | payment()
```

```
def apply() = inApply.get() >x> quote(x) >y> Email(x.addr,y) >>  
    apply()
```

```
def join() = inJoin.get() >(id,p)> validate(id,p) >>  
    ( add_client(id,p) >> Email(id.addr,welcome)  
      | renew(id)  
    ) >>  
    join()
```

```
def payment() = inPayment.get() >(id,p)> validate(id,p) >>  
    update_client(id,p) >>  
    payment()
```

Laws of Kleene Algebra

(Zero and $|$)

$$f | 0 = f$$

(Commutativity of $|$)

$$f | g = g | f$$

(Associativity of $|$)

$$(f | g) | h = f | (g | h)$$

(Idempotence of $|$)

$$f | f = f$$

(Associativity of \gg)

$$(f \gg g) \gg h = f \gg (g \gg h)$$

(Left zero of \gg)

$$0 \gg f = 0$$

(Right zero of \gg)

$$f \gg 0 = 0$$

(Left unit of \gg)

$$\text{Signal} \gg f = f$$

(Right unit of \gg)

$$f \gg x \text{ let}(x) = f$$

(Left Distributivity of \gg over $|$)

$$f \gg (g | h) = (f \gg g) | (f \gg h)$$

(Right Distributivity of \gg over $|$)

$$(f | g) \gg h = (f \gg h) | (g \gg h)$$

Laws which do not hold

(Idempotence of $|$)

$$f | f = f$$

(Right zero of \gg)

$$f \gg 0 = 0$$

(Left Distributivity of \gg over $|$)

$$f \gg (g | h) = (f \gg g) | (f \gg h)$$

Additional Laws

(Distributivity over \gg) if g is x -free

$$((f \gg g) \prec x \prec h) = (f \prec x \prec h) \gg g$$

(Distributivity over $|$) if g is x -free

$$((f | g) \prec x \prec h) = (f \prec x \prec h) | g$$

(Distributivity over \ll) if g is y -free

$$\begin{aligned} & ((f \prec x \prec g) \prec y \prec h) \\ = & ((f \prec y \prec h) \prec x \prec g) \end{aligned}$$

(Elimination of where) if f is x -free, for site M

$$(f \prec x \prec M) = f | (M \gg 0)$$

Shortest path problem

- Directed graph; non-negative weights on edges.
- Find shortest path from source to sink.

We calculate just the length of the shortest path.

Algorithm with Lights and Mirrors

- Source node sends rays of light to each neighbor.
- Edge weight is the time for the ray to traverse the edge.
- When a node receives its first ray, sends rays to all neighbors. Ignores subsequent rays.
- Shortest path length = time for sink to receive its first ray.

Expressions and Sites needed for Shortest path

succ(*u*): Publish all (v, d) , where edge (u, v) has weight d .

write(*u*, *t*): Write value t for node u . If already written, block.

read(*u*): Return value for node u . If unwritten, block.

First Algorithm

def $eval1(u, t) =$ $write(u, t) \gg$
 $Succ(u) \triangleright (v, d) \triangleright$
 $Rtimer(d) \gg$
 $eval1(v, t + d)$

Goal : $eval1(source, 0) \mid read(sink)$

First call to $eval1(u, t)$:

- The relative time in the evaluation is t .
- Length of the shortest path from source to u is t .
- $eval1$ does not publish.

Current Status

- A prototype implementation; robust, non-optimized.
- An extensive site library.
- Several small to medium applications coded.

Where we are heading

- Coding large distributed applications.
- Transactions
- Logical time
- Secure workflow.
- Adaptive workflow.

See <http://orc.csres.utexas.edu>