

Orc User Guide

Orc User Guide

Table of Contents

1. The Orc Programming Language	1
1.1. Introduction	1
1.2. Cor: A Functional Subset	1
1.2.1. Constants	2
1.2.2. Operators	2
1.2.3. Conditionals	3
1.2.4. Variables	4
1.2.5. Data Structures	4
1.2.6. Patterns	5
1.2.7. Functions	6
1.2.8. Comments	9
1.3. Orc: Orchestrating services	10
1.3.1. Communicating with external services	10
1.3.2. The concurrency combinators of Orc	12
1.3.3. Revisiting Cor expressions	14
1.3.4. Time	16
1.4. Advanced Features of Orc	16
1.4.1. Special call forms	17
1.4.2. Extensions to pattern matching	18
1.4.3. New forms of declarations	19
1.4.4. The otherwise combinator	19
2. Programming Methodology	21
2.1. Syntactic and Stylistic Conventions	21
2.1.1. Parallel combinator	21
2.1.2. Sequential combinator	21
2.1.3. Asymmetric combinator	22
2.1.4. Declarations	23
2.2. Programming Idioms	23
2.2.1. Fork-join	23
2.2.2. Parallel Or	24
2.2.3. Arbitration	25
2.2.4. Timeout	25
2.2.5. Priority	26
2.2.6. Repeat	26
2.2.7. Priority Poll	26
2.2.8. Buffer Transfer	26
2.2.9. Parallel Matching	27
2.3. Larger Examples	27
2.3.1. Dining Philosophers	27
A. Complete Syntax of Orc	29
B. Standard Library	30
B.1. Overview	30
B.2. Notation	30
B.3. Reference	31
B.3.1. core.inc	31
B.3.2. data.inc	33
B.3.3. functional.inc	38
B.3.4. list.inc	41
B.3.5. text.inc	45
B.3.6. time.inc	46
B.3.7. util.inc	47

List of Tables

1.1. Syntax of the Functional Subset of Orc (Cor)	2
1.2. Operators of Cor	3
1.3. Basic Syntax of Orc	10
1.4. Advanced Syntax of Orc	17
A.1. Complete Syntax of Orc	29

Chapter 1. The Orc Programming Language

1.1. Introduction

Orc is a programming language designed to make distributed and concurrent programs simple and intuitive to write. Orc expresses orchestration, a type of structured concurrency. It emphasizes the flow of control and gives a global view of a concurrent system. Orc is well-suited for task orchestration, a form of concurrent programming with applications in workflow, business process management, and web service orchestration. Orc provides constructs to orchestrate the concurrent invocation of services while managing time-outs, priorities, and failures of services or communication. To learn more about Orc and run your own Orc programs, visit the website: <http://orc.csres.utexas.edu/> [<http://orc.csres.utexas.edu/>].

This chapter describes the Orc programming language in three steps. In Section 1.2, we discuss a small subset of Orc called Cor. Cor is a pure functional language, which has no features for concurrency, has no state, and does not communicate with external services. Cor introduces us to the parts of Orc that are most familiar from existing programming languages, such as arithmetic operations, variables, conditionals, and functions.

Next, in Section 1.3, we consider Orc itself, which in addition to Cor, comprises external services and combinators for concurrent orchestration of those services. We show how Orc interacts with these external services, how the combinators can be used to build up complex orchestrations from simple base expressions, and how the functional constructs of Cor take on new, subtler behaviors in the concurrent context of Orc.

Finally, in Section 1.4, we discuss some additional features of Orc that extend the basic syntax. These are useful for creating large-scale Orc programs, but they are not essential to the understanding of the language.

1.2. Cor: A Functional Subset

In this section we introduce Cor, a pure functional subset of the Orc language. Users of functional programming languages such as Haskell and ML will already be familiar with many of the key concepts of Cor.

A Cor program is called an *expression*. Cor expressions are built up recursively from smaller expressions. Cor *evaluates* an expression to reduce it to some simple *value* which cannot be evaluated further, for example a list of numbers or a Boolean truth value. This value is called the *result* of the expression.

In the following subsections we introduce the concepts of Cor. First, we talk about simple constants, such as numbers and truth values, and the operations that we can perform on those values. Then we introduce conditionals (**if then else**). Then we introduce variables and binding, as a way to give a name to the value of an expression. After that, we talk about constructing data structures, and examining those structures using patterns. Lastly, we introduce functions.

The following figure describes the syntax of Cor. Each part of the syntax is explained in a subsequent section.

Table 1.1. Syntax of the Functional Subset of Orc (Cor)

$E ::=$	<i>Expression</i>
C	<i>constant value</i>
E op E	<i>operator</i>
X	<i>variable</i>
if E then E else E	<i>conditional</i>
X(E , ... , E)	<i>function call</i>
(E , ... , E)	<i>tuple</i>
[E , ... , E]	<i>list</i>
D E	<i>scoped declaration</i>
C ::= true false integer string	<i>Constant</i>
X ::= identifier	<i>Variable</i>
D ::=	<i>Declaration</i>
val P = E	<i>value declaration</i>
def X(P , ... , P) = E	<i>function declaration</i>
P ::=	<i>Pattern</i>
X	<i>variable</i>
C	<i>constant</i>
_	<i>wildcard</i>
(P , ... , P)	<i>tuple</i>
[P , ... , P]	<i>list</i>

1.2.1. Constants

The simplest expression one can write is a constant. It evaluates trivially to that constant value.

Cor has three types of constants, and thus for the moment three types of values:

- Numbers: 5, -1, 2.71828, ...
- Booleans: true and false
- Strings: "orc", "ceci n'est pas une |"

1.2.2. Operators

Cor has a standard set of arithmetic, logical, and comparison operators. As in most other programming languages, they are written in the usual infix style. They have Java-like operator precedence, which can be overridden by adding parentheses.

Examples

- 1+2 evaluates to 3.
- (98+2)*17 evaluates to 1700.
- 4 = 20 / 5 evaluates to true.

- `3-5 >= 5-3` evaluates to `false`.
- `true && (false || true)` evaluates to `true`.
- `"leap" + "frog"` evaluates to `"leapfrog"`.

Here is the full set of operators that Cor supports:

Table 1.2. Operators of Cor

Arithmetic	Comparison	Logical	String
<code>+</code> addition	<code>=</code> equality	<code>&&</code> logical and	<code>+</code> concatenation
<code>-</code> subtraction	<code>/=</code> inequality	<code> </code> logical or	
<code>*</code> multiplication	<code><</code> less than	<code>~</code> logical not	
<code>/</code> division	<code>></code> greater than		
<code>%</code> modulus	<code><=</code> less than or equal		
<code>**</code> exponent	<code>>=</code> greater than or equal		

The `=` operator can compare values of any type.

Examples

- `10 = true` evaluates to `false`.

Sometimes, there are situations where an expression is stuck, because it is attempting to perform some impossible operation and cannot reach a value. If this is the case, we say that the expression is *silent*. An expression is also silent if it depends on the result of a silent subexpression.

Examples

- `10 / 0` is silent.
- `6 + false` is silent.
- `3 + 1/0` is silent.
- `4 + true = 5` is silent.

Note that Cor is a dynamically typed language. It does not statically check the type correctness of an expression; instead, an expression with a type error is simply silent when evaluated.

1.2.3. Conditionals

A conditional expression in Cor is of the form **if** *E* **then** *F* **else** *G*. Its meaning is similar to that in other languages: the value of the expression is the value of *F* if and only if *E* evaluates to true, or the value of *G* if and only if *E* evaluates to false. Note that *G* is not evaluated at all if *E* evaluates to true, and similarly *F* is not evaluated at all if *E* evaluates to false. Thus, for example, evaluation of **if** `3 = 3` **then** `5` **else** `1/0` does not cause any error.

Unlike other languages, expressions in Cor may be silent. If *E* is silent, then the entire expression is silent. And if *E* evaluates to true but *F* is silent (or if *E* evaluates to false and *G* is silent) then the expression is silent.

The behavior of conditionals is summarized by the following table (*v* denotes a value).

E	F	G	if E then F else G
true	v	-	v
true	silent	-	silent
false	-	v	v
false	-	silent	silent
silent	-	-	silent

Examples

- `if true then 4 else 5` evaluates to 4.
- `if 2 < 3 && 5 < 4 then "blue" else "green"` evaluates to "green".
- `if true || "fish" then "yes" else "no"` is silent.
- `if false || false then 4+5 else 4>true` is silent.
- `if 0 < 5 then 0/5 else 5/0` evaluates to 0.

1.2.4. Variables

A *variable* names the value of some expression so that we can use it later. Expression `x`, where `x` is some variable name, evaluates to the result *bound* to the variable `x`.

Variables are bound using a *declaration*. A declaration is a statement that has no value of its own but instead binds one or more variables. The simplest form of declaration is `val`, which evaluates an expression and binds its result to a variable. Declarations follow the rules of lexical scoping [http://en.wikipedia.org/wiki/Lexical_scope].

```
val x = 1 + 2 + 3 + 4 + 5
val y = x + x
```

These declarations bind variable `x` to 15 and variable `y` to 30.

If the expression on the right side of a `val` is silent, then the variable is not bound, but evaluation of other declarations and expressions can continue. If an evaluated expression depends on that variable, that expression is silent.

```
val x = 1/0
val y = 4+5
if false then x else y
```

Evaluation of the declaration `val y = 4+5` and the expression `if false then x else y` may continue even though `x` is not bound. The expression evaluates to 9.

1.2.5. Data Structures

Cor supports two basic data structures, *tuples* and *lists*.

A *tuple expression* is a comma-separated sequence of at least two expressions, enclosed by parentheses. Each expression is evaluated; the value of the whole tuple expression is a tuple containing each of these values in order. If any of the expressions is silent, then the whole tuple expression is silent.

Examples

- `(1+2, 7)` evaluates to `(3,7)`.
- `("true" + "false", true || false, true && false)` evaluates to `("truefalse", true, false)`.
- `(2/2, 2/1, 2/0)` is silent.

A *list expression* is a comma-separated sequence of expressions enclosed by square brackets. It may be of any length, including zero. Each expression is evaluated; the value of the whole list expression is a list containing each of these values in order. If any of the expressions is silent, then the whole list expression is silent.

Examples

- `[1, 2+3]` evaluates to `[1, 5]`.
- `[true && true]` evaluates to `[true]`.
- `[]` evaluates trivially to `[]`, the empty list.
- `[5, 5 + true, 5]` is silent.

There is also a concatenation (*cons*) operation on lists, written `F:G`, where `F` and `G` are expressions. Its result is a new list whose first element is the value of `F` and whose remaining elements are the list value of `G`.

Examples

- `(1+3):[2+5, 6]` evaluates to `[4, 7, 6]`.
- `2:2:5:[]` evaluates to `[2, 2, 5]`.
- Suppose `t` is bound to `[3,5]`. Then `1:t` evaluates to `[1, 3, 5]`.
- `2:3` is silent, because `3` is not a list.

1.2.6. Patterns

We have seen how to construct data structures. But how do we examine them, and use them? We use *patterns*.

A pattern is a powerful way to bind variables. When writing **val** declarations, instead of just binding one variable, we can replace the variable name with a more complex pattern that follows the structure of the value, and matches its components. A pattern's structure is called its *shape*; a pattern may take the shape of any structured value. A pattern can hierarchically match a value, going deep into its structure. It can also bind an entire structure to a variable.

Examples

- **val** `(x,y) = (2+3, 2*3)` binds `x` to 5 and `y` to 6.
- **val** `[a,b,c] = ["one", "two", "three"]` binds `a` to "one", `b` to "two", and `c` to "three".
- **val** `((a,b),c) = ((1, true), (2, false))` binds `a` to 1, `b` to `true`, and `c` to `(2, false)`.

Patterns are *linear*; that is, a pattern may mention a variable name at most once. For example, (x, y, x) is not a valid pattern.

Note that a pattern may fail to match a value, if it does not have the same shape as that value. In a **val** declaration, this has the same effect as evaluating a silent expression. No variable in the pattern is bound, and if any one of those variables is later evaluated, it is silent.

It is often useful to ignore parts of the value that are not relevant. We can use the wildcard pattern, written `_`, to do this; it matches any shape and binds no variables.

Examples

- **val** $(x, _, _) = (1, (2, 2), [3, 3, 3])$ binds x to 1.
- **val** $[[_, x], [_, y]] = [[1, 3], [2, 4]]$ binds x to 3 and y to 4.

1.2.7. Functions

Like most other programming languages, Cor has the capability to define *functions*, which are expressions that have a defined name, and have some number of parameters. Functions are declared using the keyword **def**, in the following way.

```
def Add(x, y) = x+y
```

The expression on the right of the `=` is called the *body* of the function.

After defining the function, we can *call* it. A call looks just like the left side of the declaration except that the variable names (the *formal parameters*) have been replaced by expressions (the *actual parameters*).

To evaluate a call, we treat it as a sequence of **val** declarations associating the formal parameters with the actual parameters, followed by the body of the function.

```
{- Evaluation of Add(1+2, 3+4) -}  
val x = 1+2  
val y = 3+4  
x+y
```

Examples

- `Add(10, 10*10)` evaluates to 110.
- `Add(Add(5, 3), 5)` evaluates to 13.

Notice that the evaluation strategy of functions allows to call to proceed even if some of the actual parameters to be silent, so long as the values of those actual parameters are not used in the evaluation of the body.

```
def cond(b, x, y) = if b then x else y  
cond(true, 3, 5/0)
```

This evaluates to 3 even though `5/0` is silent, because `y` is not needed.

A definition or a call may have zero arguments, in which case we write `()` for the arguments.

```
def Zero() = 0
```

1.2.7.1. Recursion

Definitions can be recursive; that is, the name of a definition may be used in its own body.

```
def Sumto(n) = if n < 1 then 0 else n + Sumto(n-1)
```

Then, `Sumto(5)` evaluates to 15.

Mutual recursion is also supported.

```
def Even(n) =  
  if (n > 0) then Odd(n-1)  
  else if (n < 0) then Odd(n+1)  
  else true  
def Odd(n) =  
  if (n > 0) then Even(n-1)  
  else if (n < 0) then Even(n+1)  
  else false
```

There is no special keyword for mutual recursion; any contiguous sequence of function declarations is assumed to be mutually recursive.

1.2.7.2. Closures

Functions are actually values, just like any other value. Defining a function creates a special value called a *closure*; the name of the function is a variable and its bound value is the closure. Thus, a closure can be put into a data structure, or bound to some other variable, just like any other value.

```
def a(x) = x-3  
def b(y) = y*4  
val funs = (a,b)
```

Like any other value, a closure can be passed as an argument to another function. This means that Cor supports *higher-order* functions.

```
def Onetwosum(f) = f(1) + f(2)  
def Triple(x) = x * 3
```

Then, `Onetwosum(Triple)` is `Triple(1) + Triple(2)`, which is `1 * 3 + 2 * 3`, which evaluates to 9.

Since all declarations (including function declarations) in Cor are lexically scoped, these closures are *lexical closures*. This means that when a closure is created, if the body of the function contains any variables other than the formal parameters, the bindings for those variables are stored in the closure. Then, when the closure is called, the evaluation of the function body uses those stored variable bindings.

1.2.7.3. Lambda

Sometimes one would like to create a closure directly, without bothering to give it a name. There is a special keyword **lambda** for this purpose. By writing a function definition with the keyword **lambda** instead of a function name, that definition becomes an expression which evaluates to a closure.

```
def Onetwosum(f) = f(1) + f(2)

Onetwosum( lambda(x) = x * 3 )
{-
  identical to:
  def triple(x) = x * 3
  onetwosum(triple)
-}
```

Then, `Onetwosum(lambda(x) = x * 3)` evaluates to 9.

1.2.7.4. Clauses

The combination of functions and pattern matching offers a powerful capability: *clausal* definition of functions. We can define expressions which execute different code depending on the structure of their arguments.

Here's an example.

```
def Sum([]) = 0
def Sum(h:t) = h + Sum(t)
```

`Sum(L)` publishes the sum of the numbers in the list `L`. It has two clauses: one which matches the empty list, and one which matches any nonempty list. If its argument is an empty list, it returns 0, the appropriate sum for an empty list. If the argument is a nonempty list, it adds the first element of that list to the sum of all of the other elements. In this way, it recursively finds the sum of the list.

A function may have multiple clauses, each of which has a sequence of patterns to match each argument, and a body expression. Naturally, all clauses of a function must have the same number of arguments. Any contiguous sequence of definitions with the same name and different arguments is interpreted as a clausal definition, where each individual declaration is a clause of the larger function.

When the function is called, the clauses are tried in the order in which they appear until a match is found. If no clause matches, the call remains silent.

We allow a new form of pattern which is very useful in clausal definition of functions: a constant pattern. A constant pattern is a match only for the same constant value. We can use this to define the "base case" of a recursive function in a straightforward way.

```
{- Fibonacci numbers -}
def Fib(0) = 1
def Fib(1) = 1
def Fib(n) = if (n < 0) then 0 else Fib(n-1) + Fib(n-2)
```

This definition of the Fibonacci function is straightforward, but slow, due to the repeated work in recursive calls to `Fib`. We can define a linear-time version, again with the help of pattern matching:

```
{- Alternate definition of the Fibonacci function -}

{- A helper function: find the pair (Fibonacci(n-1), Fibonacci(n)) -}
```

```
def H(0) = (1,1)
def H(n) = H(n-1) >(x,y)> (y,x+y)

def Fib(n) = if (n < 0) then 0 else H(n) >(x,_)> x

{- Take up to the first n elements from a list -}
def Take(0,_) = []
def Take(_,[]) = []
def Take(n,h:t) = h:(Take(n-1,t))
```

Mutual recursion and clausal definitions are allowed to occur together. For example, this function takes a list and evaluates to a new list with every other element repeated:

```
def Stutter([]) = []
def Stutter(h:t) = h:h:Mutter(t)
def Mutter([]) = []
def Mutter(h:t) = h:Stutter(t)
```

`Stutter([1,2,3])` evaluates to `[1,1,2,3,3]`.

Clauses of mutually recursive functions may also be interleaved, to make them easier to read.

```
def Even(0) = true
def Odd(0) = false
def Even(n) = Odd(if n > 0 then n-1 else n+1)
def Odd(n) = Even(if n > 0 then n-1 else n+1)
```

1.2.8. Comments

Cor has two kinds of comments.

A line which begins with two dashes (--), preceded only by whitespace, is a single line comment. The region from the two dashes to the next encountered newline, inclusive, is ignored.

```
-- This is a single line comment.
-- This is also a single line comment.
```

Multiline comments are enclosed by matching braces of the form `{- -}`. Multiline comments may be nested. They may appear anywhere, even in the middle of an expression.

```
{-
  This is a
  multiline comment.
-}

{- Multiline comments {- can be nested -} -}

{- They may appear anywhere, -}
1 + {- even in the middle of an expression. -} 2 + 3
```

1.3. Orc: Orchestrating services

Cor is a pure declarative language. It has no state, since variables are bound at most once and cannot be reassigned. Evaluation of an expression results in at most one value. It cannot communicate with the outside world except by producing a value. Clearly, the full Orc language must transcend these limitations, because the orchestration of external services is critical to Orc's purpose.

As in Cor, an Orc program is an *expression*; Orc expressions are built up recursively from smaller expressions. Orc is a superset of Cor, i.e., all Cor expressions are also Orc expressions. Orc expressions are *executed*, rather than evaluated; an execution may call external services and *publish* some number of values (possibly zero). Different executions of the same expression may have completely different behaviors; they may call different services, may receive different responses from the same site, and may publish different values. Expressions in the functional subset, though, will display the same behavior in all executions.

If an execution of an expression publishes no value, it is called *silent*. If no execution of an expression can possibly publish a value, then that expression is silent. However, unlike silent expressions in Cor, a silent expression in Orc might perform some meaningful computation or communication; silence does not necessarily indicate an error.

In the following sections we discuss the features of Orc. First, we discuss how Orc communicates with external services. Then we introduce Orc's concurrency *combinators*, which combine expressions into larger orchestrations and manage concurrent executions. We have already discussed the functional subset of Orc in our coverage of Cor, but we reprise some of those topics; some Cor constructs exhibit new behaviors in the concurrent, stateful context of Orc.

The following figure summarizes the syntax of Orc as an extension of the syntax of Cor. The original Cor grammar rules are abbreviated by ellipses (...).

Table 1.3. Basic Syntax of Orc

D ::= ...	<i>Declaration</i>
site X = <i>address</i>	<i>site declaration</i>
C ::= ...	<i>Constant</i>
signal	<i>signal value</i>
E ::= ...	<i>Expression</i>
E E	<i>parallel combinator</i>
E >P> E	<i>sequential combinator</i>
E <P< E	<i>asymmetric combinator</i>
stop	<i>silent expression</i>

1.3.1. Communicating with external services

We introduce the term *site* to denote an external service which Orc may call. Sites are called using the same syntax as a function call, but with a slightly different meaning. Sites are introduced and bound to variables by a special declaration.

1.3.1.1. Calling a site

Suppose that the variable `Google` is bound to a site which invokes the Google search engine service in "I'm Feeling Lucky" mode. A call to `Google` looks just like a function call. Calling `Google` requests the URL of the top result for the given search term.

```
{- Get the top search result for "computation orchestration" -}  
Google("computation orchestration")
```

Once the Google search service determines the top result, it sends a response. The site call then publishes that response. Note that the service might never respond: Google's servers might be down, the network might be down, or the search might yield no result URL. If the service fails to respond, the site call remains silent.

A site call sends only a single request to an external service and receives at most one response. These restrictions have two important consequences. First, all of the information needed for the call must be present before contacting the service. Thus, site calls are strict; all arguments must be bound before the call can proceed. If any argument is silent, the call never occurs. Second, a site call publishes at most one value, since at most one response is received.

A site is a value, just like an integer or a list. It may be bound to a variable, passed as an argument, or published by an execution, just like any other value.

```
{-  
  Create a search site from a search engine URL,  
  bind the variable Search to that site,  
  then use that site to search for a term.  
-}  
val Search = SearchEngine("http://www.google.com/")  
Search("first class value")
```

A site is sometimes called purely for its effect on the external world; its return value is irrelevant. Many sites which do not need to return a meaningful value will instead return a *signal*: a special value which carries no information (analogous to the unit value `()` in ML). The signal value can be written as **signal** within Orc programs.

```
{-  
  Use the 'println' site to print a string, followed by  
  a newline, to an output console.  
  The return value of this site call is a signal.  
-}  
println("Hello, World!")
```

1.3.1.2. Declaring a site

A **site** declaration makes some service available as a site and binds it to a variable. The service might be an object in the host language (e.g. a class instance in Java), or an external service on the web which is accessed through some protocol like SOAP or REST, or even a primitive operation like addition. Many useful sites are already defined in the Orc standard library, documented in Appendix B. For now, we present a simple type of site declaration: using an object in the host language as a site.

The following example instantiates a Java object to be used as a site in an Orc program, specifically a Java object which provides an asynchronous buffer service. The declaration uses a fully qualified Java class name to find and load a class, and creates an instance of that class to provide the service.

```
{- Define the Buffer site -}  
site Buffer = orc.lib.state.Buffer
```

1.3.2. The concurrency combinators of Orc

Orc has three *combinators*: parallel, sequential, and asymmetric. A combinator forms an expression from two component expressions. Each combinator captures a different aspect of concurrency.

1.3.2.1. The parallel combinator

Orc's simplest combinator is `|`, the parallel combinator. Orc executes the expression `F | G`, where `F` and `G` are Orc expressions, by executing `F` and `G` concurrently. Whenever `F` or `G` communicates with a service or publishes a value, the whole execution does so as well.

```
-- Publish 1 and 2 in parallel
1 | 1+1

{-
  Access two search sites, Google and Yahoo, in parallel.

  Publish any results they return.

  Since each call may publish a value, the expression
  may publish up to two values.
-}
Google("cupcake") | Yahoo("cupcake")
```

The parallel combinator is fully associative: `(F | G) | H` and `F | (G | H)` and `F | G | H` are all equivalent.

It is also commutative: `F | G` is equivalent to `G | F`.

```
-- Publish 1, 2, and 3 in parallel
1+1 | 1 | 1+2
```

1.3.2.2. The sequential combinator

Now that we have expressions which publish multiple values, what can we do with those publications? The sequential combinator, written `F >x> G`, combines the expression `F`, which may publish some values, with another expression `G`, which will use the values as they are published; `x` transmits the values from `F` to `G`.

The execution of `F >x> G` starts by executing `F`. Whenever `F` publishes a value, a new copy of `G` is executed in parallel with `F` (and with any previous copies of `G`); in that copy of `G`, variable `x` is bound to the value published by `F`. Values published by copies of `G` are published by the whole expression, but the values published by `F` are not published by the whole expression; they are consumed by the variable binding.

```
-- Publish 1 and 2 in parallel
(0 | 1) >n> n+1

-- Publish 3 and 4 in parallel
```



```
2 >n> (n+1 | n+2)
```

```
-- Publish 0, 1, 2 and 3 in parallel
(0 | 2) >n> (n | n+1)
```

```
-- Prepend the string "Result: " to each published search result
-- The cat site concatenates any number of arguments into one string
Google("cupcake") >s> cat("Result: ", s)
| Yahoo("cupcake") >s> cat("Result: ", s)
```

The sequential combinator may be written as $F \text{ >P> } G$, where P is a pattern instead of just a variable name. Any value published by F is matched against the pattern P . If this match is successful, a new copy of G is started with all of the bindings from the match. Otherwise, the published value is simply ignored, and no new copy of G is executed.

```
-- Publish 3, 6, and 9 in arbitrary order.
(3,6,9) >(x,y,z)> ( x | y | z )
```

```
-- Filter out values of the form (_,false)
-- Publishes 4 and 6
( (4,true) | (5,false) | (6,true) ) >(x,true)> x
```

We may also omit the variable entirely, writing >> . This is equivalent to using a wildcard pattern: >_>

We may want to execute an expression just for its effects, and hide all of its publications. We can do this using >> together with the special expression **stop**, which is always silent.

```
{-
  Print two strings to the console,
  but don't publish the return values of the calls.
-}
( println("goodbye") | println("world") ) >> stop
```

The sequential combinator is right associative: $F \text{ >x> } G \text{ >y> } H$ is equivalent to $F \text{ >x> } (G \text{ >y> } H)$. It has higher precedence than the parallel combinator: $F \text{ >x> } G \text{ | } H$ is equivalent to $(F \text{ >x> } G) \text{ | } H$.

The right associativity of the sequential combinator makes it easy to bind variables in sequence and use them together.

```
{-
  Publish the cross product of {1,2} and {3,4}:
  (1,3), (1,4), (2,3), and (2,4).
-}
(1 | 2) >x> (3 | 4) >y> (x,y)
```

1.3.2.3. The asymmetric combinator

The asymmetric combinator, written $F \text{ <x< } G$, allows us to block a computation waiting for a result, or terminate a computation. The execution of $F \text{ <x< } G$ starts by executing F and G in parallel. Whenever F publishes a value, that value is published by the entire execution. When G publishes its first value, that

value is bound to x in F , and then the execution of G is immediately *terminated*. A terminated expression cannot call any sites or publish any values.

During the execution of F , any part of the execution that depends on x will be suspended until x is bound (to the first value published by G). If G never publishes a value, that part of the execution is suspended forever.

```
-- Publish either 5 or 6, but not both
x+2 <x< (3 | 4)
```

```
-- Query Google and Yahoo for a search result
-- Print out the result that arrives first; ignore the other result
println(result) <result< ( Google("cupcake") | Yahoo("cupcake") )
```

Though a terminated execution may not make any new calls, the calls that it has already made will continue normally; their responses are simply ignored. This may have surprising consequences when a call has side effects, as in the following example.

```
{-
  This example actually prints both "true" and "false" to the
  console, regardless of which call responds first.
-}
stop <x< println("true") | println("false")
```

Both of the `println` calls are initiated before either one of them publishes a value and terminates the expression. Once the expression is terminated, no new calls occur, but the other `println` call still proceeds and still has the effect of printing its message to the console.

The asymmetric combinator may include a full pattern P instead of just a variable name. Any value published by G is matched against the pattern P . If this match is successful, then G terminates and all of the bindings of the pattern P are made in F . Otherwise, the published value is simply ignored and G continues to execute.

```
-- Publish either 9 or 25, but not 16.
x*x <(x,true)< ( (3,true) | (4,false) | (5,true) )
```

Note that even if $(4, false)$ is published before $(3, true)$ or $(5, true)$, it is ignored. The right side continues to execute and will publish one of $(3, true)$ or $(5, true)$.

The asymmetric combinator is left associative: $F <x< G <y< H$ is equivalent to $(F <x< G) <y< H$. It has lower precedence than the parallel combinator: $F <x< G | H$ is equivalent to $F <x< (G | H)$.

1.3.3. Revisiting Cor expressions

Some Cor expressions have new behaviors in the context of Orc, due to the introduction of concurrency and of sites.

1.3.3.1. `val`

The declaration `val x = G`, followed by expression F , is actually just a prefix form of the asymmetric combinator $F <x< G$. Thus, `val` shares all of the behavior of the asymmetric combinator, which we have already described. (This is also true when a pattern is used instead of variable name x).

1.3.3.2. **if**

The conditional expression **if** *E* **then** *F* **else** *G* is actually a derived form based on a site named **if**. The **if** site takes a boolean argument; it returns a signal if that argument is `true`, or remains silent if the argument is `false`.

if *E* **then** *F* **else** *G* is equivalent to $(\text{if}(b) \gg F \mid \text{if}(\sim b) \gg G) \leq b < E$.

When the **else** branch of a conditional is unneeded, we can use the **if** site and the sequential combinator to write the simpler expression **if**(*E*) \gg *F*.

1.3.3.3. Treating Orc expressions functionally

We let Orc expressions with concurrency combinators occur as subexpressions of Cor expressions, for example as operands to arithmetic operations, or as arguments to function calls. The execution of an Orc expression may publish many values and perform many site calls. What does such an expression mean in a purely functional context, where only one value is expected? For example, what does $2 + (3 \mid 4)$ publish?

The specific syntactic contexts we are interested in are as follows (*E* is any Orc expression):

<i>E op E</i>	<i>operand</i>
if <i>E</i> then ...	<i>conditional test</i>
<i>X</i> (... , <i>E</i> , ...)	<i>call argument</i>
(... , <i>E</i> , ...)	<i>tuple element</i>
[... , <i>E</i> , ...]	<i>list element</i>

These are the contexts where only one value is expected from *E*. Whenever an Orc expression appears in such a context, it executes until it publishes its first value, and then it is terminated. The published value is used in the context as if it were the result of evaluating the expression.

```
-- Publish either 5 or 6
2 + (3 | 4)
```

```
-- Publish exactly one of 0, 1, 2 or 3
(0 | 2) + (0 | 1)
```

To be precise, whenever an Orc expression appears in a Cor context, it is treated as if it was on the right side of an asymmetric combinator, using a fresh variable name to fill in the hole. Thus, $C[E]$ (where *E* is an Orc expression and *C* is a Cor context) is equivalent to the expression $C[x] \leq x < E$.

1.3.3.4. Functions

The body of a function in Orc may be any Orc expression; thus, function bodies in Orc are executed rather than evaluated, and may engage in communication and publish multiple values.

A function call in Orc, as in Cor, binds the values of its actual parameters to its formal parameters, and then executes the function body with those bindings. Whenever the function body publishes a value, the function call publishes that value. Thus, unlike a site call, or a pure functional Cor call, an Orc function call may publish many values.

```
-- Publish all integers in the interval 1..n, in arbitrary order.
def range(n) = if (n > 0) then (n | range(n-1)) else stop

-- Publish 1, 2, and 3 in arbitrary order.
range(3)
```

In the context of Orc, function calls are not strict. When a function call executes, it begins to execute the function body immediately, and also executes the argument expressions in parallel. When an argument expression publishes a value, it is terminated, and the corresponding formal parameter is bound to that value in the execution of the function body. Whenever an execution in the function body uses a formal parameter that has not been bound yet, that execution suspends until that parameter is bound to a value.

1.3.4. Time

Orc is designed to communicate with the external world, and one of the most important characteristics of the external world is the passage of time. Orc implicitly interacts with the passage of time by calling external services which take time to respond. However, Orc can also explicitly wait for some amount of time, using the special site `Rtimer`.

The site `Rtimer` is a relative timer. It takes as an argument an amount of time to wait. It waits for exactly that amount of time, and then responds with a signal. By default, this argument is in milliseconds.

```
-- Print "red", wait for 3 seconds (3000 ms), and then print "green"
println("red") >> Rtimer(3000) >> println("green") >> stop
```

The following example defines a metronome, which publishes a signal once per second, indefinitely.

```
def metronome() = signal | Rtimer(1000) >> metronome()
```

We can also use `Rtimer` together with the asymmetric combinator to enforce a timeout.

```
{-
  Publish the result of a Google search.
  If it takes more than 5 seconds, time out.
-}
result
  <result< ( Google("impatience")
             | Rtimer(5000) >> "Search timed out.")
```

We present many more examples of programming techniques using real time in Chapter 2.

1.4. Advanced Features of Orc

In this section we introduce some advanced features of Orc. These include curried function definitions and curried calls, writing an arbitrary expression as the target of a call, a special syntax for writing calls in an object-oriented style, extensions to pattern matching, new forms of declarations, and a new concurrency combinator.

The following figure summarizes further extensions to the syntax of Orc. The Orc and Cor grammar rules presented earlier are abbreviated by ellipses (...). The item `G+` means "one or more instances of `G` concatenated together".

Table 1.4. Advanced Syntax of Orc

$E ::= \dots$	<i>Expression</i>	
$E\ G^+$		<i>generalized call</i>
$E\ ;\ E$		<i>otherwise combinator</i>
$G ::=$	<i>Argument group</i>	
$(\ E\ ,\ \dots\ ,\ E\)$		<i>curried arguments</i>
$\ .\ field$		<i>field access</i>
$P ::= \dots$	<i>Pattern</i>	
$!P$		<i>publish pattern</i>
$P\ as\ X$		<i>as pattern</i>
$D ::= \dots$	<i>Declaration</i>	
class $X = classname$		<i>class declaration</i>
include " $filename$ "		<i>inclusion</i>

1.4.1. Special call forms

1.4.1.1. The `.` notation

In many object-oriented programming languages, one calls a method or accesses a field of an object using the dot operator; for example, `obj.m()` calls the method `m` of the object `obj`.

There is a special kind of site call in Orc which serves a similar purpose. One may write `x.msg`, for any identifiers `x` and `msg`. This attempts to treat the value bound to `x` as a site, and call it with a special *message* value. If the site understands the message `msg` (for example, if `x` is bound to a Java object with a field called `msg`), the site interprets the message and responds with some appropriate value. If the site does not understand the message sent to it, it does not respond, and no publication occurs. If the value cannot be interpreted as a site, no call is made.

Typically this capability is used so that sites may be syntactically treated like objects, with multiple methods and fields. For example, a channel `c` might understand the messages `get` and `put`, to get values from and put values on that channel, respectively. Such calls would be written `c.get()`, or `c.put(6)`.

A call such as `c.put(6)` actually occurs in two steps. First `c.put` sends the message `put` to the site `c`; this publishes a site whose only purpose is to put values on the channel. Next, that site is called on the argument `6`, sending `6` on the channel. Readers familiar with functional programming will recognize this technique as *currying*.

1.4.1.2. Currying

It is sometimes useful to *stage* the arguments to a function; that is, rather than writing a function on two arguments, one instead writes a function on one argument which returns a function taking the second argument and performing the remainder of the evaluation.

This technique is known as *currying* and it is common in functional programming languages. We can write curried functions using closures. Suppose we want to define a curried addition function on two arguments, and later apply that function to the arguments 3 and 4. We could write such a program in the following way:

```
def Add(a) = ( lambda(b) = a+b )
val f = Add(3)
f(4)
```

This defines a function `add` which, given an argument `a`, creates a function which will take an argument `b` and add `a` to `b`. It then creates the function which adds 3 to its argument, binds that to `f`, and then invokes `f` on 4 to yield `3+4`.

When defining a curried function, we have abstracted it in two steps, and when applying it we have written two separate calls. However, this is verbose and not very clear. Orc has a special syntax for curried definitions and curried applications that will simplify both of these steps. Here is the previous program, written in this simplified syntax:

```
def Add(a)(b) = a+b
Add(3)(4)
```

Definitions may have multiple argument sequences; they are enclosed in parentheses and concatenated. Curried function calls chain together multiple applications in a similar way.

1.4.2. Extensions to pattern matching

1.4.2.1. Publish pattern

A publish pattern, written `!p`, will publish the value that matches the pattern `p` if the match is successful.

```
(1,2,3) >(x,!y,!z)> stop
```

This publishes 2 and 3.

Note that a publish pattern will not publish if the overall match fails, even if its particular match succeeds:

```
((1,2,3) | (4,5,6)) >(1,!x,y)> stop
```

This publishes only 2. Even though the pattern `x` matches the value 5, the overall pattern `(1,!x,y)` does not match the value `(4,5,6)`, so 5 is not published.

1.4.2.2. As pattern

In taking apart a value with a pattern, it is often useful to capture intermediate results.

```
val (a,b) = ((1,2),(3,4))
val (ax,ay) = a
val (bx,by) = b
```

We can use the `as` keyword to simplify this program fragment, giving a name to an entire subpattern. Here is an equivalent version of the above code.

```
val ((ax,ay) as a, (bx,by) as b) = ((1,2),(3,4))
```

1.4.3. New forms of declarations

1.4.3.1. `class` declaration

When Orc is run on top of an object-oriented programming language, classes from that language may be used as sites in Orc itself, via the **class** declaration.

```
{- Use the String class from Java's standard library as a site -}  
class String = java.lang.String  
val s = String("foo")  
s.concat("bar")
```

This program binds the variable `String` to the constructor of Java's `String` class. When it is called, it constructs a new instance of `String`, passing the given arguments to the constructor.

This instance of `String` is a Java object; its methods are called and its fields are accessed using the `.` notation, just as one would expect in Java.

1.4.3.2. `include` declaration

It is often convenient to group related declarations into units that can be shared between programs. The **include** declaration offers a simple way to do this. It names a source file containing a sequence of Orc declarations; those declarations are incorporated into the program as if they had textually replaced the include declaration. An included file may itself contain **include** declarations.

```
{- Contents of fold.inc -}  
def foldl(f,[],s) = s  
def foldl(f,h:t,s) = foldl(f,t,f(h,s))  
  
def foldr(f,l,s) = foldl(f,rev(l),s)  
  
{- This is the same as inserting the contents of fold.inc here -}  
include "fold.inc"  
  
def sum(L) = foldl(lambda(a,b) = a+b, L, 0)  
  
sum([1,2,3])
```

Note that these declarations still obey the rules of lexical scope. Also, Orc does not detect shared declarations; if the same file is included twice, its declarations occur twice.

1.4.4. The `otherwise` combinator

Orc has a fourth concurrency combinator: the *otherwise* combinator, written `F ; G`. The execution of `F ; G` executes `F`, calling sites and publishing values as usual. If `F` *completes*, and has not published any values by the time it completes, then `G` executes. If `F` did publish one or more values, then `G` is ignored.

We determine when an expression completes using the following rules.

- A Cor expression completes when it is fully evaluated; if it is silent, it completes immediately.
- A site call completes when it has published a value.

- A site call completes silently if it is blocked on a variable, and the expression which would bind that variable has already completed.
- A site call completes silently if the site *refuses* the call, as explained below.
- $F \mid G$ completes when its subexpressions F and G have both completed.
- $F >x> G$ completes when F has completed and all copies of G have completed.
- $F <x< G$ completes when F has completed, and G has either completed or has been terminated after publishing a value.
- $F ; G$ completes either when F has published a value and subsequently completed, or when F and G have both completed.
- **stop** completes immediately.

A site refuses a call when it knows that it will never respond to that call. This may happen for a variety of reasons. A site may be specifically designed to never respond under certain circumstances; for example, `if(false)` will never respond, so `if` refuses that call. A site may depend on some connection that has failed in a detectable way; for example, `Google("five tons of flax")` if it is running on a computer which has no internet connection. A site may require some resource which is unavailable; for example, a site call based on a website may fail if the website itself is down or has moved (a 404 error). Generally speaking, any underlying exception or error which prevents the site call from ever continuing will trigger a refusal. Refusal is used only to implement the otherwise combinator; it is not used elsewhere in the language.

The otherwise combinator was not present in the original formulation of the Orc concurrency calculus; it has been added to support computation and iteration over strictly finite data. Sequential programs conflate the concept of producing a value with the concept of completion. Orc separates these two concepts; variable binding combinators like `>x>` and `<x<` handle values, whereas `;` detects the completion of an execution.

Chapter 2. Programming Methodology

In Chapter 1, we described the syntax and semantics of the Orc language. Now, we turn our attention to how the language is used in practice, with guidelines on style, programming tricks, and a number of common concurrency patterns.

2.1. Syntactic and Stylistic Conventions

In this section we suggest some syntactic conventions for writing Orc programs. None of these conventions are required by the parser; newlines are used only to disambiguate certain corner cases in parsing, and other whitespace is ignored. However, following programming conventions helps to improve the readability of programs, so that programmer intent is more readily apparent.

2.1.1. Parallel combinator

When the combined expressions are small, write them all on one line.

```
F | G | H
```

Note that we do not need parentheses here, since `|` is fully associative and commutative.

When the combined expressions are large enough to take up a full line, write one expression per line, with the subsequent expressions preceded by `|`. Indent the first expression to improve readability.

```
    long expression  
| long expression  
| long expression
```

A sequence of parallel expressions often form the left hand side of a sequential combinator. Since the sequential combinator has higher precedence, use parentheses to group the combined parallel expressions together.

```
( expression  
| expression  
) >x>  
another expression
```

2.1.2. Sequential combinator

When the combined expressions are small, write a cascade of sequential combinators all on the same line.

```
F >x> G >y> H
```

When the combined expressions are large enough to take up a full line, write one expression per line; each line ends with the combinator which binds the publications produced by that line.

```
long expression >x>  
long expression >y>  
long expression
```

For very long-running expressions, or expressions that span multiple lines, write the combinators on separate lines, indented, between each expression.

```
very long expression
  >x>
very long expression
  >y>
very long expression
```

2.1.3. Asymmetric combinator

When the combined expressions are small, write them on the same line:

```
F <x< G
```

When multiple asymmetric combinators are used to bind multiple variables (especially when the scoped expression is long), start each line with a combinator, aligned and indented, and continue with the expression.

```
long expression
  <x< G
  <y< H
```

The asymmetric combinator is not often written in its explicit form in Orc programs. Instead, the **val** declaration is often more convenient, since it is semantically equivalent and mentions the variable **x** before its use in scope, rather than after.

```
val x = G
val y = H
long expression
```

Additionally, when the variable is used in only one place, and the expression is small, it is often easier to use a nested expression. For example,

```
val x = G
val y = H
M(x, y)
```

is equivalent to

```
M(G, H)
```

Sometimes, we use the asymmetric combinator simply for its capability to terminate expressions and get a single publication; binding a variable is irrelevant. This is a special case of nested expressions. We use the identity site **let** to put the expression in the context of a function call.

For example,

```
x <x< F | G | H
```

is equivalent to

```
let (F | G | H)
```

The translation uses an asymmetric combinator, but we don't need to write the combinator, name an irrelevant variable, or worry about precedence (since the expression is enclosed in parentheses as part of the call).

2.1.4. Declarations

Declarations should always end with a newline.

```
def add(x,y) = x + y
val z = 7
add(z,z)
```

While the parser does not require a newline to end a declaration, it uses the newline to disambiguate certain corner cases in parsing, such as function application.

When the body of a declaration spans multiple lines, start the body on a new line after the = symbol, and indent the entire body.

```
def f(x,y) =
  declaration
  declaration
  body expression
```

2.2. Programming Idioms

In this section we look at some common idioms used in the design of Orc programs. Many of these idioms will be familiar to programmers using concurrency, and they are very simple to express in Orc.

2.2.1. Fork-join

One of the most common concurrent idioms is a *fork-join*: run two processes concurrently, and wait for a result from each one. This is very easy to express in Orc. Whenever we write a **val** declaration, the process computing that value runs in parallel with the rest of the program. So if we write two **val** declarations, and then form a tuple of their results, this performs a fork-join.

```
val x = F
val y = G
(x,y)
```

Fork-joins are a fundamental part of all Orc programs, since they are created by all nested expression translations. In fact, the fork-join we wrote above could be expressed even more simply as just:

```
(F,G)
```

2.2.1.1. Example: Machine initialization

In Orc programs, we often use fork-join and recursion together to dispatch many tasks in parallel and wait for all of them to complete. Suppose we want to initialize a list of machines; for each machine *m*,

calling `m.init()` initializes `m` and then publishes a signal when initialization is complete. The function `initAll` does exactly that:

```
def initAll([]) = signal
def initAll(m:ms) = ( m.init() , initAll(ms) ) >> signal
```

For each machine, we fork-join the initialization of that machine (`m.init()`) with the initialization of the remaining machines (`initAll(ms)`). Thus, all of the initializations proceed in parallel, and the function returns a signal only when every machine in the list has completed its initialization.

Note that if some machine fails to initialize, and does not return a signal, then the initialization procedure will never complete. We will see later how to better manage such failures, using a timeout or a quorum.

2.2.1.2. Example: Simple parallel auction

We can also use a recursive fork-join to obtain a value, rather than just signaling completion. Suppose we have a list of bidders in a sealed-bid, single-round auction. Calling `b.ask()` requests a bid from the bidder `b`. We want to ask for one bid from each bidder, and then return the highest bid. The function `auction` performs such an auction (Max finds the maximum of its arguments):

```
def auction([]) = 0
def auction(b:bs) = Max(b.ask(), auction(bs))
```

2.2.1.3. Barrier synchronization

Suppose we have an expression of the following form, where `F` and `G` are expressions and `M` and `N` are sites:

$$M() \triangleright x \triangleright F \mid N() \triangleright y \triangleright G$$

However, we would also like to *synchronize* the completion of `M` and `N`, so that neither `F` nor `G` starts executing until both `M` and `N` have published. This is a particular example of a fork-join:

$$(M() , N()) \triangleright (x,y) \triangleright (F \mid G)$$

We assume, naturally, that `x` does not occur free in `G`, nor `y` in `F`.

2.2.2. Parallel Or

Next we consider a classic example of parallel programming which builds up from fork-join: a parallel-or program. Given two expressions `F` and `G` which may publish boolean values (or might stay silent forever), we want to find the disjunction of their results as soon as possible. In particular, we know that if one of the expressions publishes `true`, then the disjunction must be `true`; it is not necessary to wait for the other expression to publish a value.

Here is the code:

```
val a = F
val b = G
val r = if(a) >> true
```

```

      | if(b) >> true
      | (a || b)
r

```

Recall that the `||` operator is strict; if one of `a` or `b` is not bound, it cannot compute a result. So, we introduce two extra branches of the form `if(a) >> true`, which wait in parallel for each variable to become `true` and then publish the result `true`. That way, the parallel-or can evaluate to `true` based on only one of the result, even if the other result is not forthcoming.

2.2.3. Arbitration

Another important concurrent idiom is *arbitration*: choosing between two tasks based on the output of another task. This is simple to accomplish in Orc using conditionals. This program executes an expression `P`, which produces a boolean value, and then uses that value to choose whether to subsequently run `F` or `G`.

```

val b = P
if b then F else G

```

2.2.4. Timeout

One of the most powerful idioms available in Orc is a *timeout*: execute an expression for at most a specified amount of time. We accomplish this using `val` and a timer. The following programs runs `F` for at most one second to get a value from it. If it does not publish within one second, the value defaults to 0.

```

let( F | Rtimer(1000) >> 0 )

```

2.2.4.1. Auction with timeout

In our earlier auction example, the auction may never complete if one of the bidders does not respond. We can add a timeout so that a bidder has at most 8 seconds to provide a bid:

```

def auction([]) = 0
def auction(b:bs) =
  val bid = b.ask() | Rtimer(8000) >> 0
  Max(bid, auction(bs))

```

Now, the auction is guaranteed to complete in at most 8 seconds.

2.2.4.2. Detecting timeout

Sometimes, rather than just yielding a default value, we would like to determine whether an expression has timed out, and if so, perform some other computation. To detect the timeout, we pair the expression with `true` and the timer with `false`, so if the expression does time out, then we can detect it using the truth value. We then test the truth value, performing the usual computation if no timeout occurred, or some error-correcting computation if the expression did time out.

Here, we run the expression `F` with a time limit `t`. If it publishes within the time limit, we execute `G` (the result of `F` is bound to `r`). Otherwise, we execute `H`.

```

val (r,b) = (F,true) | (Rtimer(t),false)

```

```
if b then G else H
```

2.2.5. Priority

We can also use a timer to give a window of priority to one computation over another. In this example, we run expressions *F* and *G* concurrently. *F* has priority for the first second: if *F* publishes, its value is used immediately, but if *G* publishes, that publication is suspended for the first second and may be superseded by a publication from *F* during that time. If neither *F* nor *G* publishes within a second, then whichever publishes first after that point is the winner.

```
val x = F
val y = G
let( y | Rtimer(1000) >> x )
```

2.2.6. Repeat

Recall the definition of *metronome* from the previous chapter:

```
def metronome() = signal | Rtimer(1000) >> metronome()
```

We can use *metronome* and the sequential combinator to repeat an expression, so that copies of that expression will execute at regular intervals.

```
{-
  Publish "tick" once every second.
  Publish "tock" once every second, with an initial 500ms delay.
  The publications alternate: tick tock tick tock ...
-}
metronome() >> "tick"
| Rtimer(500) >> metronome() >> "tock"
```

2.2.7. Priority Poll

Suppose we have a list of buffers. We would like to poll these buffers periodically to see if they have any available data. Furthermore, this list of buffers is ordered by priority; the first buffer in the list has the highest priority, so it should be polled first, and if it has no data, then the next buffer should be polled, and so on.

Here is a function which polls a prioritized list of buffers in this way. It publishes the first item that it finds, removing it from the originating buffer. If all buffers are empty, it remains silent. We use the *getnb* ("get non-blocking") method of the buffer, which retrieves the first available item if there is one, or else remains silent and completes immediately if the buffer is empty (it does not wait for an item to become available).

```
def PriorityPoll([]) = stop
def PriorityPoll(b:bs) = b.getnb() ; PriorityPoll(bs)
```

2.2.8. Buffer Transfer

Suppose we have two buffers, and would like to empty one buffer and put its contents into the other buffer. We use the following function to do so.

```
def Append(x,y) = y.put(x.getnb()) >> Append(x,y) ; signal
```

Append takes the elements from `x` one at a time (using `getnb`) and recurses. When `x` is empty, `getnb` completes and is silent, so the left side of the `;` completes and the function publishes a signal, indicating that the transfer is finished.

2.2.9. Parallel Matching

Matching a value against multiple patterns, as we have seen it so far, is a linear process, and requires a `def` whose clauses have patterns in their argument lists. Such a match is linear; each pattern is tried in order until one succeeds.

What if we want to match a value against multiple patterns in parallel, executing every clause that succeeds? Fortunately, this is very easy to do in Orc. Suppose we have an expression `F` which publishes pairs of integers, and we want to publish a signal for each 3 that occurs. We write:

```
F >x>
  ( x >(3,_)> signal
  | x >(_,3)> signal )
```

The interesting case is the pair `(3,3)`, which is counted twice because both patterns match it in parallel.

This is a useful idiom even if the patterns are mutually exclusive, to avoid creating a helper function just to handle the pattern match.

2.3. Larger Examples

2.3.1. Dining Philosophers

```
def shuffle(a,b) = if (random(2) = 1) then (a,b) else (b,a)

def take(a,b) =
  a.acquire() >> b.obtain() ;
  a.release() >> take(shuffle(a,b))

def drop(a,b) = (a.release(), b.release()) >> signal

def phil(a,b) =
  def thinking() = Rtimer(random(1000)) >> hungry()
  def hungry() = take(a,b) >> eating()
  def eating() = Rtimer(random(1000)) >> drop(a,b) >> thinking()
  thinking()

def dining(n) =
  val forks = makeArray(n, Semaphore)
  def phils(0) = stop
  def phils(i) = phil(forks(i%n), forks(i-1)) | phils(i-1)
  phils(n)
```

dining(5)

Appendix A. Complete Syntax of Orc

Table A.1. Complete Syntax of Orc

$E ::=$	<i>Expression</i>	
C		<i>constant value</i>
$ E \text{ op } E$		<i>operator</i>
$ X$		<i>variable</i>
$ \text{if } E \text{ then } E \text{ else } E$		<i>conditional</i>
$ E G^+$		<i>call</i>
$ (E , \dots , E)$		<i>tuple</i>
$ [E , \dots , E]$		<i>list</i>
$ D E$		<i>scoped declaration</i>
$ E \mathbf{!} E$		<i>parallel combinator</i>
$ E >P> E$		<i>sequential combinator</i>
$ E <P< E$		<i>asymmetric combinator</i>
$ E ; E$		<i>otherwise combinator</i>
$ \text{stop}$		<i>silent expression</i>
$G ::=$	<i>Argument group</i>	
(E , \dots , E)		<i>arguments</i>
$. \text{field}$		<i>field access</i>
$C ::= \text{true} \text{false} \text{integer} \text{string} \text{signal}$	<i>Constant</i>	
$X ::= \text{identifier}$	<i>Variable</i>	
$D ::=$	<i>Declaration</i>	
$\text{val } P = E$		<i>value declaration</i>
$ \text{def } X(P , \dots , P) = E$		<i>function declaration</i>
$ \text{site } X = \text{address}$		<i>site declaration</i>
$ \text{class } X = \text{classname}$		<i>class declaration</i>
$ \text{include " filename "}$		<i>inclusion</i>
$P ::=$	<i>Pattern</i>	
X		<i>variable</i>
$ C$		<i>constant</i>
$ _$		<i>wildcard</i>
$ (P , \dots , P)$		<i>tuple</i>
$ [P , \dots , P]$		<i>list</i>
$!P$		<i>publish pattern</i>
$ P \text{ as } X$		<i>as pattern</i>

Appendix B. Standard Library

B.1. Overview

The standard library is a set of declarations implicitly available to all Orc programs. In this section we give an informal description of the standard library, including the type of each declaration and a short explanation of its use.

Orc programs are expected to rely on the host language and environment for all but the most essential sites. For example, in the Java implementation of Orc, the entire Java standard library is available to Orc programs via **class** declarations. Therefore the Orc standard library aims only to provide convenience for the most common Orc idioms, not the complete set of features needed for general-purpose programming.

The standard library is divided into groups of related declarations:

- `core.inc`: Essential primitive sites and operators.
- `data.inc`: General concurrent and sequential data structures.
- `functional.inc`: Higher-order functional programming idioms.
- `list.inc`: Operations on lists.
- `text.inc`: Operations on strings.
- `time.inc`: Real and logical time.
- `util.inc`: Miscellaneous utility functions.

B.2. Notation

Type signatures are written using a combination of Java- and Haskell-like syntax. Example: **site** `let<A>(A) :: A`. The signature consists of a declaration keyword followed by a name, type parameters, formal parameter types, and finally the return type.

Declaration keywords include **site**, **def**, **val**, and **pattern**. The first three keywords should be familiar. The last, **pattern**, indicates a binding for a site which can be used in a pattern match: the site must take one argument (the value to be matched against), and return a tuple (the values to bind to pattern variables). For example, **pattern** `M(Integer) :: (Integer, Integer)` would be used like `3 >M(x,y)> x+y`.

Object members (methods and fields) are each declared separately. The name of a member is written in the form `Type.member`, e.g. `Foo.get` refers to the `get` member of an object of type `Foo`. The object type can include type parameters which may be referenced by the member type.

A list with element type `A` is written `[A]`. A tuple with element types `A`, `B`, ... is written `(A, B, ...)`. The Top and Bottom types are written `Top` and `Bot` respectively. The type of **signal** is `Signal`. The meanings of other named types should be clear from context.

A function type (when given as an argument or return type) is written with the keyword **lambda** followed by the argument types and return type, e.g. **lambda** `(Integer) :: Integer`. Multiple formal parameter types are syntactic sugar for currying, e.g. **site** `Foo()() :: Signal` is equivalent to **site** `Foo() :: lambda () :: Signal`.

B.3. Reference

B.3.1. core.inc

<code>let</code>	When applied to a single argument, return that argument (behaving as the identity function).
<code>let</code>	When applied to zero, two, or more arguments, return the arguments in a tuple.
<code>if</code>	Fail silently if the argument is false.
<code>error</code>	Fail with the given error message.
<code>(+)</code>	Add two numbers.
<code>(-)</code>	Subtract the right argument from the left.
<code>(0-)</code>	Return the additive inverse of the argument.
<code>(*)</code>	Multiply two numbers.
<code>(**)</code>	<code>a ** b</code> computes <code>a</code> to the <code>b</code> th power.
<code>(/)</code>	If both arguments have integral types, perform integral division, rounding towards zero.
<code>(%)</code>	<code>a % b</code> computes the remainder of <code>a / b</code> , satisfying the equation <code>a = (a/b) * b + (a%b)</code> .
<code>(<)</code>	Is the left argument less than the right?
<code>(<=)</code>	Is the left argument less than or equal to the right?
<code>(>)</code>	Is the left argument greater than the right?
<code>(>=)</code>	Is the left argument greater than or equal to the right?
<code>(=)</code>	Is the left argument equivalent to the right?
<code>(/=)</code>	Is it not the case that the left argument is equivalent to the right, as described for <code>(=)</code> ?
<code>(~)</code>	Compute the logical negation of the argument.
<code>(&&)</code>	Compute the logical conjunction of the arguments.
<code>()</code>	Compute the logical disjunction of the arguments.
<code>(:)</code>	Prepend an element to a list.
<code>(:)</code>	Deconstruct a list into a head and tail.

B.3.1.1. `site let<A>(A) :: A`

When applied to a single argument, return that argument (behaving as the identity function).

B.3.1.2. `site let<A, ...>(A, ...) :: (A, ...)`

When applied to zero, two, or more arguments, return the arguments in a tuple.

B.3.1.3. `site if(Boolean) :: Signal`

Fail silently if the argument is false.

```
    if(false) >> "Never publishes"
  | if(true) >> "Always publishes"
```

B.3.1.4. site error(String) :: Bot

Fail with the given error message.

```
def assert(b) =
  if b then signal else error("assertion failed")
assert(false)
```

B.3.1.5. site (+)(Number, Number) :: Number

Add two numbers.

B.3.1.6. site (-)(Number, Number) :: Number

Subtract the right argument from the left.

B.3.1.7. site (0-)(Number) :: Number

Return the additive inverse of the argument. This site is written without the 0 when it appears as an operator, i.e. $-a$

B.3.1.8. site (*)(Number, Number) :: Number

Multiply two numbers.

B.3.1.9. site ()(Number, Number) :: Number**

$a ** b$ computes a to the b th power.

B.3.1.10. site (/)(Number, Number) :: Number

If both arguments have integral types, perform integral division, rounding towards zero. Otherwise, perform floating-point division. For example:

```
    7 / 3 -- publishes 2
  | 7 / 3.0 -- publishes 2.333...
```

B.3.1.11. site %(Number, Number) :: Number

$a \% b$ computes the remainder of a / b , satisfying the equation $a = (a/b)*b+(a\%b)$. For a full description, see the Java Language Specification, 3rd edition [http://java.sun.com/docs/books/jls/third_edition/html/expressions.html#15.17.3].

B.3.1.12. site (<)(Comparable, Comparable) :: Boolean

Is the left argument less than the right?

B.3.1.13. site (<=)(Comparable, Comparable) :: Boolean

Is the left argument less than or equal to the right?

B.3.1.14. `site (>)(Comparable, Comparable) :: Boolean`

Is the left argument greater than the right?

B.3.1.15. `site (>=)(Comparable, Comparable) :: Boolean`

Is the left argument greater than or equal to the right?

B.3.1.16. `site (=)(Top, Top) :: Boolean`

Is the left argument equivalent to the right? Values are considered equivalent if one can be substituted locally for another without affecting the behavior of the program. Therefore two different immutable objects are equivalent if they share equivalent content, while two different mutable objects are never equivalent.

When `a` and `b` are arbitrary Java objects, `a = b` is interpreted as `a.equals(b)`, which may not follow the rules given above.

B.3.1.17. `site (/=)(Top, Top) :: Boolean`

Is it not the case that the left argument is equivalent to the right, as described for `(=)`?

B.3.1.18. `site (~)(Boolean) :: Boolean`

Compute the logical negation of the argument.

B.3.1.19. `site (&&)(Boolean, Boolean) :: Boolean`

Compute the logical conjunction of the arguments. This is not a short-circuiting operator; both arguments must be evaluated and available before the result is computed.

B.3.1.20. `site (||)(Boolean, Boolean) :: Boolean`

Compute the logical disjunction of the arguments. This is not a short-circuiting operator; both arguments must be evaluated and available before the result is computed.

B.3.1.21. `site (:)<A>(A, [A]) :: [A]`

Prepend an element to a list.

B.3.1.22. `pattern (:)<A>([A]) :: (A, [A])`

Deconstruct a list into a head and tail.

B.3.2. `data.inc`

`Semaphore`

Construct a semaphore with the given number of available items.

`Semaphore.acquire`

Acquire an item.

`Semaphore.release`

Release an item.

Buffer	Create a new buffer (FIFO channel) of unlimited size.
Buffer.get	Get an item from the buffer.
Buffer.getnb	Get an item from the buffer.
Buffer.put	Put an item in the buffer.
Buffer.close	Close the buffer.
SyncChannel	Create a synchronous channel.
SyncChannel.get	Receive an item over the channel.
SyncChannel.put	Send an item over the channel.
Cell	Create a write-once cell.
Cell.read	Read a value from the cell.
Cell.write	Write a value to the cell.
Ref	Create an empty rewritable reference.
Ref	Create a rewritable reference containing the provided value.
Ref.read	Read the value of the reference.
Ref.write	Write a value to the reference.
Null	Return a Java null value.
Array	Create a new array of the given size.
Array	Create a new primitive array of the given size with the given primitive type.
Array.get	Get the element of the array given by the index, counting from 0.
Array.set	Set the element of the array given by the index, counting from 0.
Array.slice	Return a copy of the portion of the array with indices covered by the given half-open range.
Array.length	Return the size of the array.
Some	Construct an available optional value.
Some	Deconstruct an available optional value.
None	Construct an unavailable optional value.
None	Deconstruct an unavailable optional value.
Left	Construct a "left" member of a union.
Left	Deconstruct a "left" member of a union.
Right	Construct a "right" member of a union.
Right	Deconstruct a "right" member of a union.
fst	Return the first element of a pair.
snd	Return the second element of a pair.
swap	Swap the elements of a pair.

B.3.2.1. site Semaphore(Integer) :: Semaphore

Construct a semaphore with the given number of available items.

```
val lock = Semaphore(1)
lock.acquire() >>
println("Entering critical section") >>
println("Leaving critical section") >>
lock.release() >> stop
```

B.3.2.2. site Semaphore.acquire() :: Signal

Acquire an item. If none are available, block until one becomes available.

B.3.2.3. site Semaphore.release() :: Signal

Release an item. This may be called to add items beyond the initial number of available items with which the semaphore was constructed.

B.3.2.4. site Buffer<A>() :: Buffer<A>

Create a new buffer (FIFO channel) of unlimited size.

```
val b = Buffer()
Rtimer(1000) >> b.put(10) >> stop
| b.get()
```

B.3.2.5. site Buffer<A>.get() :: A

Get an item from the buffer. If no items are available, block until one becomes available.

B.3.2.6. site Buffer<A>.getnb() :: A

Get an item from the buffer. If no items are available, fail silently.

B.3.2.7. site Buffer<A>.put(A) :: Signal

Put an item in the buffer.

B.3.2.8. site Buffer<A>.close() :: Signal

Close the buffer. This has the effect of immediately causing any blocked calls to get to fail silently. In addition, any subsequent calls to get or put will fail silently.

B.3.2.9. site SyncChannel<A>() :: SyncChannel<A>

Create a synchronous channel.

```
val c = SyncChannel()
c.put(10) | Rtimer(1000) >> c.get()
```

B.3.2.10. site SyncChannel<A>.get() :: A

Receive an item over the channel. If no sender is available, block until one becomes available.

B.3.2.11. site SyncChannel<A>.put(A) :: Signal

Send an item over the channel. If no receiver is available, block until one becomes available.

B.3.2.12. site Cell<A>() :: Cell<A>

Create a write-once cell.

```
val c = Cell()
c.write(5)
| c.write(10)
| c.read()
```

B.3.2.13. site Cell<A>.read() :: A

Read a value from the cell. If the cell does not yet have a value, block until it receives one.

B.3.2.14. site Cell<A>.write() :: Signal

Write a value to the cell. If the cell already has a value, fail silently.

B.3.2.15. site Ref<A>() :: Ref<A>

Create an empty rewritable reference.

```
val r = Ref()
Rtimer(1000) >> r.write(5) >> stop
| println(r.read()) >>
  r.write(10) >>
  println(r.read()) >>
  stop
```

B.3.2.16. site Ref<A>(A) :: Ref<A>

Create a rewritable reference containing the provided value.

B.3.2.17. site Ref<A>.read() :: A

Read the value of the reference. If the reference is empty, block until it is no longer empty.

B.3.2.18. site Ref<A>.write() :: Signal

Write a value to the reference.

B.3.2.19. site Null() :: Bot

Return a Java null value. This is only necessary to interface with certain Java libraries, Orc programs should use the `None()` constructor instead of null values.

B.3.2.20. site Array<A>(Integer) :: Array<A>

Create a new array of the given size.


```
def fillArray(a, f) =  
  for(0, a.length()) >i>  
    a.set(i, f(i)) >>  
  stop  
  ; a  
val a = fillArray(Array(3), let)  
a.get(0) | a.get(1) | a.get(2)
```

B.3.2.21. site Array<A>(Integer, String) :: Array<A>

Create a new primitive array of the given size with the given primitive type. The primitive type should match the element type of the array, although a typechecker may not be able to verify this. This constructor is only necessary when interfacing with certain Java libraries; most programs will just use the `Array<A>(Integer)` constructor.

B.3.2.22. site Array<A>.get(Integer) :: A

Get the element of the array given by the index, counting from 0.

B.3.2.23. site Array<A>.set(Integer, A) :: Signal

Set the element of the array given by the index, counting from 0.

B.3.2.24. site Array<A>.slice(Integer, Integer) :: Array<A>

Return a copy of the portion of the array with indices covered by the given half-open range. The result array is still indexed counting from 0.

B.3.2.25. site Array<A>.length() :: Integer

Return the size of the array.

B.3.2.26. site Some<A>(A) :: Option<A>

Construct an available optional value.

B.3.2.27. pattern Some<A>(Option<A>) :: (A)

Deconstruct an available optional value.

```
Some(3) >Some(x) > x
```

B.3.2.28. site None<A>(A) :: Option<A>

Construct an unavailable optional value.

B.3.2.29. pattern None<A>(Option<A>) :: ()

Deconstruct an unavailable optional value.

```
None() >None()> true
| Some(3) >None()> false
```

B.3.2.30. site Left<A>(A) :: Either<A>

Construct a "left" member of a union.

B.3.2.31. pattern Left<A>(Either<A>) :: A

Deconstruct a "left" member of a union.

```
Left(3) >x> (
  x >Right(_)> "right"
| x >Left(_)> "left"
)
```

B.3.2.32. site Right<A>(A) :: Either<A>

Construct a "right" member of a union.

B.3.2.33. pattern Right<A>(Either<A>) :: A

Deconstruct a "right" member of a union.

```
Right(3) >x> (
  x >Right(_)> "right"
| x >Left(_)> "left"
)
```

B.3.2.34. def fst<A,B>(A,B) :: A

Return the first element of a pair.

B.3.2.35. def snd<A,B>(A,B) :: B

Return the second element of a pair.

B.3.2.36. def swap<A,B>((A,B)) :: (B,A)

Swap the elements of a pair.

B.3.3. functional.inc

apply
curry
curry3
uncurry
uncurry3

Apply a function to a list of arguments.
Curry a function of two arguments.
Curry a function of three arguments.
Uncurry a function of two arguments.
Uncurry a function of three arguments.

flip	Flip the order of parameters of a two-argument function.
const	Create a function which returns a constant value.
delay	Return a thunk which applies a function to a single argument.
ignore	From a function of no arguments, create a function of one argument, which is ignored.
compose	Compose two single-argument functions.
until	Iterate a function until a predicate returns false.
while	Iterate a function while a predicate returns true.
repeat	Call a function sequentially, publishing each value returned by the function until it no longer returns.
fork	Call a list of functions in parallel, publishing all values published by the functions.
sequence	Call a list of functions in sequence, publishing a signal whenever the last function publishes.
join	Call a list of functions in parallel and publish a signal once all functions have completed.
por	Parallel or.
pand	Parallel and.
collect	Run a function, collecting all publications in a list.

B.3.3.1. site apply<A,B>(lambda (A, ...) :: B, [A]) :: B

Apply a function to a list of arguments.

B.3.3.2. def curry<A,B,C>(lambda (A,B) :: C)(A)(B) :: C

Curry a function of two arguments.

B.3.3.3. def curry3<A,B,C,D>(lambda (A,B,C) :: D)(A)(B)(C) :: D

Curry a function of three arguments.

B.3.3.4. def uncurry<A,B,C>(lambda (A)(B) :: C)(A, B) :: C

Uncurry a function of two arguments.

B.3.3.5. def uncurry3<A,B,C,D>(lambda (A)(B)(C) :: D)(A,B,C) :: D

Uncurry a function of three arguments.

B.3.3.6. def flip<A,B,C>(lambda (A, B) :: C)(B, A) :: C

Flip the order of parameters of a two-argument function.

B.3.3.7. def const<A>(A)() :: A

Create a function which returns a constant value.

B.3.3.8. def delay<A,B>(lambda (A) :: B, A)() :: B

Return a thunk which applies a function to a single argument.

B.3.3.9. def ignore<A,B>(lambda () :: B)(A) :: B

From a function of no arguments, create a function of one argument, which is ignored.

B.3.3.10. def compose<A,B,C>(lambda (B) :: C, lambda (A) :: B)(A) :: C

Compose two single-argument functions.

B.3.3.11. def until<A>(lambda (A) :: Boolean, lambda (A) :: A)(A)

Iterate a function until a predicate returns false. Publishes every value returned by the function.

```
until(  
  lambda (n) = (n = 5),  
  lambda (n) = n+1 )(0)
```

B.3.3.12. def while<A>(lambda (A) :: Boolean, lambda (A) :: A)(A)

Iterate a function while a predicate returns true. Publishes every value passed to the function.

```
while(  
  lambda (n) = (n <= 5),  
  lambda (n) = n+1 )(0)
```

B.3.3.13. def repeat<A>(lambda () :: A) :: A

Call a function sequentially, publishing each value returned by the function until it no longer returns.

B.3.3.14. def fork<A>([lambda () :: A]) :: A

Call a list of functions in parallel, publishing all values published by the functions.

B.3.3.15. def sequence<A>([lambda () :: A]) :: Signal

Call a list of functions in sequence, publishing a signal whenever the last function publishes.

B.3.3.16. def join<A>([lambda () :: A]) :: Signal

Call a list of functions in parallel and publish a signal once all functions have completed.

B.3.3.17. `def por(lambda () :: Boolean, lambda () :: Boolean) :: Boolean`

Parallel or. Evaluate two boolean functions in parallel, publishing a value as soon as possible, and terminating any unnecessary ongoing computation.

B.3.3.18. `def pand(lambda () :: Boolean, lambda () :: Boolean) :: Boolean`

Parallel and. Evaluate two boolean functions in parallel, publishing a value as soon as possible, and terminating any unnecessary ongoing computation.

B.3.3.19. `def collect<A>(lambda () :: A) :: [A]`

Run a function, collecting all publications in a list.

```
collect(delay(signals, 5))
```

B.3.4. `list.inc`

<code>each</code>	Publish every value in a sequence, simultaneously.
<code>map</code>	Apply a function to every element of a sequence (in parallel), producing a list of the results.
<code>reverse</code>	Reverse a list.
<code>filter</code>	Return a list containing only those elements which satisfy the predicate.
<code>head</code>	Return the first element of a sequence.
<code>tail</code>	Return all but the first element of a sequence.
<code>init</code>	Return all but the last element of a sequence.
<code>last</code>	Return the last element of a sequence.
<code>empty</code>	Is the sequence empty?
<code>index</code>	Return the <code>n</code> th element of a list, counting from 0.
<code>append</code>	Concatenate two lists.
<code>foldl</code>	Reduce a list using the given left-associative binary operation and initial value.
<code>foldl1</code>	A special case of <code>foldl</code> which uses the last element of the list as the initial value.
<code>foldr</code>	Reduce a list using the given right-associative binary operation and initial value.
<code>foldr1</code>	A special case of <code>foldr</code> which uses the last element of the list as the initial value.
<code>zip</code>	Combine two lists into a list of pairs.
<code>unzip</code>	Split a list of pairs into a pair of lists.
<code>length</code>	Return the number of elements in a list.

take	Return the first <i>n</i> elements of a list.
drop	Return the elements of a list after the first <i>n</i> .
member	Is the given item a member of the list?
merge	Merge two sorted lists.
mergeBy	Merge two lists using the given less-than relation.
sort	Sort a list.
sortBy	Sort a list using the given less-than relation.
group	Given a list of pairs, group together the second elements of consecutive pairs with equal first elements.
groupBy	Given a list of pairs, group together the second elements of consecutive pairs with equal first elements, using the given equality relation.
range	Generate a list of integers in the given half-open range.
any	Do any the elements of the list match the predicate?
all	Do all the elements of the list match the predicate?
sum	Compute the sum of all numbers in a list.

B.3.4.1. def each<A>([A]) :: A

Publish every value in a sequence, simultaneously.

B.3.4.2. def map<A,B>(lambda (A) :: B, [A]) :: [B]

Apply a function to every element of a sequence (in parallel), producing a list of the results.

B.3.4.3. def reverse<A>([A]) :: [A]

Reverse a list.

B.3.4.4. def filter<A>(lambda (A) :: Boolean, [A]) :: [A]

Return a list containing only those elements which satisfy the predicate.

B.3.4.5. def head<A>([A]) :: A

Return the first element of a sequence.

B.3.4.6. def tail<A>([A]) :: [A]

Return all but the first element of a sequence.

B.3.4.7. def init<A>([A]) :: [A]

Return all but the last element of a sequence.

B.3.4.8. def last<A>([A]) :: A

Return the last element of a sequence.

B.3.4.9. def empty<A>([A]) :: Boolean

Is the sequence empty?

B.3.4.10. def index<A>(Integer, [A]) :: A

Return the nth element of a list, counting from 0.

B.3.4.11. def append<A>([A], [A]) :: [A]

Concatenate two lists.

B.3.4.12. def foldl<A,B>(lambda (B, A) :: B, B, [A]) :: B

Reduce a list using the given left-associative binary operation and initial value. Given the list [x1, x2, x3, ...] and initial value x0, returns f(... f(f(f(x0, x1), x2), x3) ...)

```
foldl(flip((:)), [], [1,2,3])
```

B.3.4.13. def foldl1<A>(lambda (A, A) :: A, [A]) :: A

A special case of foldl which uses the last element of the list as the initial value.

B.3.4.14. def foldr<A,B>(lambda (A, B) :: B, B, [A]) :: B

Reduce a list using the given right-associative binary operation and initial value. Given the list [..., x3, x2, x1] and initial value x0, returns f(... f(x3, f(x2, f(x1, x0))) ...)

```
foldr((+), 0, [1,2,3])
```

B.3.4.15. def foldr1<A>(lambda (A, A) :: A, [A]) :: A

A special case of foldr which uses the last element of the list as the initial value.

B.3.4.16. def zip<A,B>([A], [B]) :: (A, B)

Combine two lists into a list of pairs. The length of the shortest list determines the length of the result.

B.3.4.17. def unzip<A,B>([(A,B)]) :: ([A], [B])

Split a list of pairs into a pair of lists.

B.3.4.18. def length<A>([A]) :: Integer

Return the number of elements in a list.

B.3.4.19. `def take<A>(Integer, [A]) :: [A]`

Return the first *n* elements of a list.

B.3.4.20. `def drop<A>(Integer, [A]) :: [A]`

Return the elements of a list after the first *n*.

B.3.4.21. `def member<A>(A, [A]) :: Boolean`

Is the given item a member of the list?

B.3.4.22. `def merge<A extends Comparable>([A], [A]) :: [A]`

Merge two sorted lists.

```
merge([1,2,3], [2,4,5])
```

B.3.4.23. `def mergeBy<A>(lambda (A,A) :: Boolean, [A], [A]) :: [A]`

Merge two lists using the given less-than relation.

B.3.4.24. `def sort<A extends Comparable>([A]) :: [A]`

Sort a list.

```
sort([1,3,2])
```

B.3.4.25. `def sortBy<A>(lambda (A,A) :: Boolean, [A]) :: [A]`

Sort a list using the given less-than relation.

B.3.4.26. `def group<A,B>([(A,B)]) :: [(A,[B])]`

Given a list of pairs, group together the second elements of consecutive pairs with equal first elements.

```
group([(1,1), (1,2), (2,3), (3,4), (1,3)])
```

B.3.4.27. `def groupBy<A,B>(lambda (A,A) :: Boolean, [(A,B)]) :: [(A,[B])]`

Given a list of pairs, group together the second elements of consecutive pairs with equal first elements, using the given equality relation.

B.3.4.28. `def range(Integer, Integer) :: [Integer]`

Generate a list of integers in the given half-open range.

B.3.4.29. def any<A>(lambda (A) :: Boolean, [A]) :: Boolean

Do any the elements of the list match the predicate? Uses parallel-or to return a result as soon as possible and terminate unnecessary computation.

B.3.4.30. def all<A>(lambda (A) :: Boolean, [A]) :: Boolean

Do all the elements of the list match the predicate? Uses parallel-and to return a result as soon as possible and terminate unnecessary computation.

B.3.4.31. def sum([Number]) :: Number

Compute the sum of all numbers in a list.

B.3.5. text.inc

cat	Concatenate one or more values as strings.
print	Print one or more values as strings, concatenated, to standard output.
println	Print one or more values as strings, concatenated, to standard output, with each value followed by a newline.
parseInt	Parse a string as a BigInteger.
parseBool	Parse a string as a Boolean (true/false).
lines	Split a string into lines, which are substrings terminated by an endline or the end of the string.
unlines	Append a linefeed, "\n", to each string in the sequence and concatenate the results.
words	Split a string into words, which are sequences of non-whitespace characters separated by whitespace.
unwords	Concatenate a sequence of strings with a single space between each string.

B.3.5.1. site cat(Top, ...) :: String

Concatenate one or more values as strings. For Java objects, this will call `toString()` to convert the object to a String.

B.3.5.2. site print(Top, ...) :: Signal

Print one or more values as strings, concatenated, to standard output. For Java objects, this will call `toString()` to convert the object to a String.

B.3.5.3. site println(Top, ...) :: Signal

Print one or more values as strings, concatenated, to standard output, with each value followed by a newline. For Java objects, this will call `toString()` to convert the object to a String.

B.3.5.4. site parseInt(String) :: BigInteger

Parse a string as a BigInteger.

B.3.5.5. site parseBool(String) :: Boolean

Parse a string as a Boolean (true/false).

B.3.5.6. def lines(String) :: [String]

Split a string into lines, which are substrings terminated by an newline or the end of the string. DOS, Mac, and Unix newline conventions are all accepted. Endline characters are not included in the result.

B.3.5.7. def unlines([String]) :: String

Append a linefeed, "\n", to each string in the sequence and concatenate the results.

B.3.5.8. def words(String) :: [String]

Split a string into words, which are sequences of non-whitespace characters separated by whitespace.

B.3.5.9. def unwords([String]) :: String

Concatenate a sequence of strings with a single space between each string.

B.3.6. time.inc

Rtimer	Publish a signal after the given number of milliseconds.
Clock	A call to <code>Clock</code> creates a new relative clock.
Ltimer	Publish a signal after the given number of logical timesteps.
Metronome	Publish a signal every second, indefinitely.
MetronomeT	Publish a signal at regular intervals, indefinitely.

B.3.6.1. site Rtimer(Integer) :: Signal

Publish a signal after the given number of milliseconds.

B.3.6.2. site Clock()() :: Integer

A call to `Clock` creates a new relative clock. Calling a relative clock returns the number of milliseconds which have elapsed since the clock was created.

```
val c = Clock()  
Rtimer(1000) >> c()
```

B.3.6.3. site Ltimer(Integer) :: Signal

Publish a signal after the given number of logical timesteps. A logical timestep is complete as soon as all outstanding site calls (other than calls to `Ltimer`) have published.

B.3.6.4. def Metronome() :: Signal

Publish a signal every second, indefinitely.

B.3.6.5. def MetronomeT(Integer) :: Signal

Publish a signal at regular intervals, indefinitely. The period is given by the argument, in milliseconds.

B.3.7. util.inc

random	Return a random Integer value chosen from the range of all possible 32-bit Integer values.
urandom	Returns a pseudorandom, uniformly distributed Double value between 0 and 1, inclusive.
UUID	Return a random (type 4) UUID represented as a string.
Thread	Given a site, return a new site which calls the original site in a separate thread.
signals	Publish the given number of signals, simultaneously.
for	Publish all values in the given half-open range, simultaneously.

B.3.7.1. site random() :: Integer

Return a random Integer value chosen from the range of all possible 32-bit Integer values.

site random(Integer) :: Integer Return a pseudorandom, uniformly distributed Integer value between 0 (inclusive) and the specified value (exclusive). If the argument is 0, fail silently.

B.3.7.2. site urandom() :: Double

Returns a pseudorandom, uniformly distributed Double value between 0 and 1, inclusive.

B.3.7.3. site UUID() :: String

Return a random (type 4) UUID represented as a string.

B.3.7.4. site Thread(Site) :: Site

Given a site, return a new site which calls the original site in a separate thread. This is necessary when calling a Java site which does not cooperate with Orc's scheduler and may block for an unpredictable amount of time.

A limited number of threads are reserved in a pool for use by this site, so there is a limit to the number of blocking, uncooperative sites that can be called simultaneously.

B.3.7.5. def signals(Integer) :: Signal

Publish the given number of signals, simultaneously.

```
signals(5)
```

B.3.7.6. def for(Integer, Integer) :: Integer

Publish all values in the given half-open range, simultaneously.

```
for(1,6)
```